

Three Generations of Healthcare IT: From the Digital Record to the Computable Care Process

Alexander Apartsin¹, Yehudit Aperstein²

¹School of Computer Science, Faculty of Sciences, Holon Institute of Technology (HIT), 52 Golomb St., Holon 5810201, Israel

²Intelligent Systems, Afeka Academic College of Engineering, 218 Bnei Efraim St., Tel-Aviv 6910717, Israel

Correspondence: Yehudit Aperstein, apersteiny@afeka.ac.il. Co-author: A. Apartsin, alexanderap@hit.ac.il. ORCID: Apartsin 0009-0000-7007-3529, Aperstein 0000-0001-6390-9463. Submitted to JAMIA as a Perspective.

Abstract

Objective. Healthcare IT is usually organized by the technologies it adopts. We instead organize it by the unit of information a system makes computable, and describe a computational layer whose object is patient-specific clinical intent.

Approach. We give criteria for a computational layer, derive three (record, clinical state, and a proposed layer of intent), and formalize the Actionable Clinical Record (ACR) as the third layer's atomic object.

Discussion. The framework distinguishes prescribed, observed, and intended process; existing standards represent intent once it is structured but do not recover it from natural communication, the capability we localize. The ACR is complementary to FHIR workflow resources, guidelines, and process mining; a companion feasibility study illustrates tractability for one narrow subproblem.

Conclusion. Computable clinical intent is a coherent research direction; the ACR, its readiness ladder, and an executable-correctness evaluation framework are reusable constructs for subsequent work to extend, evaluate, or falsify.

Keywords: electronic health records; clinical workflow; actionable clinical record; clinical intent; natural language processing; large language models; FHIR; care-process automation

1. Introduction

Healthcare information technology is best analyzed through the computational objects it supports rather than the technologies it adopts.^[1,2] Organized this way, the field's history is compact: the first generation made the clinical *record* computable, the second the clinical *fact*, and a third, which we propose, makes patient-specific clinical *intent* computable from communication. We use *generation* for the emergence of a new computational layer; the layers accumulate rather than replace one another.

A limitation persists after the first two layers are in place. When a clinician writes “repeat the complete blood count in two weeks” or “refer to cardiology if symptoms persist,” the instruction is decisive yet captured only as narrative text, its action, timing, condition, and owner outside the computable substrate. The cost is measurable: test-result follow-up often fails,^[3] in one large health system only about one-third of specialist referrals reached a completed visit,^[4] and such loop-closure failures, rooted in ambiguous responsibility and care-coordination breakdowns,^[5] contribute to diagnostic error and avoidable harm.^[6,7,8,9]

This Perspective offers a framework and vocabulary rather than a new empirical study: layer criteria and a three-layer organization (Section 2), the Actionable Clinical Record (ACR) as the third layer’s atomic object (Section 5), an executable-correctness evaluation framework and companion feasibility study (Section 7), and a research agenda (Section 8).

2. A Framework of Computational Layers

We distinguish a computational layer from an incremental trend by five markers: a new atomic *unit of computability*; an *external driver*; an *enabling technology*; a new *class of computation* the unit permits; and a *residual limitation* that motivates the next layer’s object. The layers are *cumulative*, each building on those beneath it (Figure 1), and differ in epistemic status: the first two are *retrospective abstractions* over a well-documented history,^[1,2] whereas the third is a *prospective hypothesis*, offered as an analytic lens rather than an established periodization. The lens is complementary to the data–information–knowledge hierarchy^[10] and to digital-maturity models such as EMRAM,^[11,12] which measure adopted capability; our concern is *which* object becomes computable (Table 1).



Figure 1. The three computational layers of healthcare IT, organized by the unit of information each makes computable. Layers accumulate rather than replace; layer 3 is proposed. Alt text: three stacked boxes, record at the base, clinical state above it, and clinical intent at the top, with an upward arrow indicating that each layer builds on the one below.

Axis	Layer 1: Record	Layer 2: Clinical state	Layer 3: Clinical intent (proposed)
Unit of computability	Clinical record	Discrete clinical fact	Intended clinical action (ACR)
Question answered	What was documented?	What is true about the patient?	What is supposed to happen?
External driver	Adoption policy and reimbursement	Interoperability and quality measurement	Care continuity and patient safety
Enabling technology	Transactional databases; messaging	Terminologies; structured data; interoperability; clinical NLP	Language models; symbolic reasoning; workflow semantics
Class of computation	Store, retrieve, exchange	Query, reason, predict	Schedule, coordinate, monitor, close
Canonical output	Document	Coded fact	Actionable Clinical Record
Residual limitation	Meaning remains in narrative	Intent remains in communication	Safe, reliable execution

Table 1. The three computational layers against the framework markers. Layers 1 and 2 are retrospective; layer 3 is proposed. The residual limitation of each layer motivates the unit of the next.

3. The First Two Layers: Record and Clinical State

The first layer digitized the clinical record and its transactions, replacing the paper chart with electronic capture, storage, retrieval, transmission, and operations such as order

entry and results routing. In the United States, HITECH and Meaningful Use moved adoption to near-universal within a decade,^[13,14] extending the older problem-oriented-record ideal.^[15] Its enabling technology, the transactional database, stored and moved information without making its clinical meaning computable across systems.

The second layer turned the record into computable clinical *state*. Its enabling stack was broad, not a single technique: structured data entry, controlled terminologies,^[16] computerized order entry, data warehouses, and interoperability standards^[17,18] mattered as much as language processing, one route by which narrative facts entered the semantic layer (MedLEE, cTAKES).^[19,20] These enabled the first broad wave of computation *on* clinical data: decision support, quality measurement, registries, and research.^[21] A tension remained: encoding meaning into structured fields captures facts, while much clinical reasoning and many intended actions remain expressed only in narrative.^[22]

4. The Remaining Computational Gap

The first two layers made clinical *state* computable, but many patient-specific *intended actions* remain non-computable while they are expressed only in unstructured communication. Healthcare IT can represent process, but only once a person or system has structured it: FHIR can represent patient-specific intent through *CarePlan*, *ServiceRequest*, and *Task*, while *PlanDefinition* defines reusable protocols that are instantiated per patient;^[17] computer-interpretable guidelines and care-pathway models encode executable protocols;^[23,24,25] and process mining reconstructs pathways from event logs.^[26,27] These paradigms are complementary and often provide the infrastructure to operationalize its outputs. The gap is isolated by distinguishing three kinds of process

information (Table 2): *prescribed* (what should generally happen), the domain of guidelines and decision support; *observed* (what actually happened), the domain of records and process mining; and *intended* (what is supposed to happen for a specific patient), often expressed in communication rather than structured directly. Existing paradigms represent and operationalize intended process once it has been structured, but none recovers it from unstructured communication; that recovery, and its conversion into executable form, is the subject of the third layer.

Process type	Meaning	Typical source	Existing paradigm
Prescribed	What should generally happen	guideline / protocol	computer-interpretable guidelines; decision support
Observed	What actually happened	event log / EHR	process mining
Intended	What is supposed to happen for a specific patient	clinical communication	represented once structured (FHIR); recovery from communication proposed here

Table 2. Three kinds of process information. The categories describe the semantic status of a process statement, not exclusive source types. The proposed third layer targets recovery of patient-specific intended process that remains expressed in natural communication.

5. Computable Clinical Intent

We propose that the third layer makes patient-specific clinical intent computable by recovering it from communication as a structured, executable representation, using one terminology throughout: clinical *communication* may express *intent*; an intent decomposes into intended *actions*; each action is an *Actionable Clinical Record*; a set of ACRs and related events is a clinical *workflow*; and the enacted sequence over time is the *care process*.

5.1 The Actionable Clinical Record

Definition 1 (Actionable Clinical Record). An Actionable Clinical Record is a source-grounded representation of a patient-specific intended clinical action, with the semantic attributes required to interpret, execute, monitor, or audit it: the tuple

118 ACR = $\langle$ action, target, actor, temporal constraint, condition, dependency, status, prove-
 119 nance, confidence $\rangle$
 120 in which *action*, *status*, and *provenance* are required; *target*, *actor*, *tem-*
 121 *poral constraint*, *condition*, and *dependency* are populated when communicated; and
 122 *confidence* records recovery-system uncertainty on inferred attributes.

123 The attributes are defined semantically (Figure 2). The *temporal constraint* is a
 124 normalized time semantics rather than a date alone, admitting absolute, relative, event-
 125 anchored, and recurring forms; *provenance* grounds the record in its source span,
 126 speaker, and encounter; and *confidence*, a calibrated per-attribute uncertainty, routes
 127 uncertain records to review. *Target* links to a layer-2 coded concept, and *dependency*
 128 encodes ordering.



130 **Figure 2.** The ACR schema. Alt text: nine labeled boxes forming the ACR tuple; action, status,
 131 and provenance are shaded to mark them as required; target, actor, temporal constraint, condition, de-
 132 pendency, and confidence are outlined as optional.

133 The third layer spans two distinct axes that together prevent over- and under-
 134 claiming. A representational *readiness* ladder runs *mentioned* $\rightarrow$ *interpreted* $\rightarrow$ *action-*
 135 *able* $\rightarrow$ *executable*: recognizing that “complete blood count” appears in a note belongs to
 136 layer 2; extracting “repeat complete blood count” is a partial step; adding “two weeks”
 137 and a responsible actor makes it actionable; resolving a due date and a workflow object
 138 makes it executable. Separately, the ACR’s *status* attribute records the operational
 139 lifecycle of the recovered action, from *recovered* and *confirmed* to *active* and then *com-*

pleted or cancelled, without changing its representational readiness; detecting completion or escalation closes the loop.

5.2 From communication to workflow

The third layer is a bridge rather than a replacement electronic health record (Figure 3).

Its inputs are the range of clinical communication (discharge instructions, conversation,

handovers, referrals, messages, calls, transcripts); its output is a set of ACRs that maps

onto infrastructure the field already has.^[17,28] The mapping situates the ACR *upstream*

of these resources rather than as a competitor, translating attributes into workflow elements

according to the target resource and profile: for example, *action* and *target* into a

Task or *ServiceRequest*, *temporal constraint* into the resource-appropriate timing element,

actor into performer or owner semantics, and *dependency* into *basedOn/partOf*

links. Recovery *provenance* can itself be represented with a FHIR *Provenance* resource

or profile, while *confidence* generally remains metadata of the recovery system. Existing

systems execute such instructions once structured; the third layer recovers them from

communication.




Figure 3. The ACR as the layer between clinical communication and existing workflow infrastructure. Alt text: communication sources on the left feed a central Actionable Clinical Record layer, which maps to existing infrastructure such as FHIR resources, schedulers, and referral and results systems on the right.

Follow-up is the most-documented case, but it is one of a broader class of communication-derived process objects sharing the ACR representation: conditional escalation, medication transitions, handoff commitments, referrals, pending-result responsibility, and patient self-management. Ambient documentation and recent shared tasks extract structured orders directly from clinical conversations,^[29,30] showing these channels carry actionable content.

5.3 Relation to existing paradigms

Two mature lines can be mistaken for the third layer. Clinical NLP already extracts concepts, relations, events, temporality, and follow-up recommendations, increasingly with language models,^[19,20,31,32,33,34,35,36] and directly extracts physician action items and medical decisions from clinical notes.^[37,38] Even order extraction from conversation targets a flat schema scored by text overlap;^[30] the distinction the ACR draws is operational, requiring that a recovered action be source-grounded, temporally normalized, and carry the actor, condition, and dependency structure needed to execute rather than only index it. Computer-interpretable guidelines encode generic, population-level protocols;^[23,24] the third layer recovers patient-specific records from a particular encounter, capturing what a clinician instructed for this patient.

6. Reliability and Architecture

A more capable generator does not remove the requirements the third layer imposes: auditability, verifiable temporal correctness, and a calibrated signal of when to defer are properties of the output contract, not of model capability. General-purpose models interpret clinical language well^[39] but can produce unreliable or unsupported outputs;^[40]

this matters most in binding an action to its time and normalizing it into an executable schedule, a setting where general-purpose temporal reasoning is known to be weak.^[41]

These requirements motivate a hybrid decomposition in which learned components interpret and explicit components enforce verifiable structure.^[42,43] Because some workflow semantics (calendar arithmetic, dependency ordering, status transitions) are formally specified, learned extraction supplies candidate actions and arguments^[44] while deterministic reasoning resolves times and due dates^[45,46,47,48] and selective prediction routes uncertain cases to a clinician. The requirements can be stated as invariants: every acted-upon record traces to a source span, speaker, and encounter;^[49] deterministic semantics are computed, not generated; each inferred attribute carries a calibrated confidence; the system may abstain; higher-consequence actions receive tighter human review;^[50] an immutable trace links communication to action for audit;^[51] and extracted content stays distinct from inferred. Human confirmation is the default until task-specific calibration and prospective safety evidence justify higher automation.

7. Evaluation Framework and Feasibility

Because the third layer acts, it should be evaluated on *executable correctness* rather than text overlap. We propose an evaluation framework for ACR recovery whose measures are reported separately: action detection; argument extraction; action–time, condition, and actor linking; temporal-normalization error; whole-record exact match; unsupported- and omitted-action rates; calibration and selective-risk behavior; and source-provenance accuracy. These assess ACR recovery; loop-closure outcome, a downstream operational endpoint shaped by implementation and workflow rather than

recovery correctness, is reported separately. This framework, not any single score, is the contribution we intend others to adopt.

A companion study offers a formative feasibility test of one minimal task, follow-up-instruction extraction, evaluating a hybrid neural-symbolic pipeline on a controlled benchmark against generative baselines.^[52] On the action–time pairing the executable representation requires, the hybrid pipeline substantially outperforms general-purpose language models, which recognize actions but bind them to times unreliably. This is an existence proof for one narrow subproblem, not evidence for the framework, whose value rests on the constructs it defines.

8. Research Agenda

The framework defines a program in four domains. **Representation:** a consensus ACR ontology, with attribute semantics, constraint types, provenance, and reconciliation when communications revise or revoke a plan. **Inference and reliability:** ACR recovery from multilingual and multimodal communication, temporal reasoning, cross-message reconciliation, calibrated abstention, distinguishing patient from clinician intent, and separating who spoke from who is responsible for acting. **Evaluation:** authentic-data benchmarks, the framework of Section 7 across tasks, and cross-institution and cross-language generalization. **Integration and impact:** ACR-to-FHIR write-back into scheduling, results-management, and referral systems,^[17,23] review-queue design, regulatory classification, and whether closed-loop operation improves outcomes.

Moving from synthetic corpora to de-identified real-world notes across languages and documentation styles is the central empirical risk, and safe deployment presupposes the invariants of Section 6. The proposal makes testable predictions: machine-

actionable use of FHIR workflow resources should grow relative to free-text follow-up; evaluation should shift toward executable-correctness measures; reliable systems should expose source-linked, auditable records with selective review; and ambient documentation should extend from notes to tasks and orders.^[29,53,54] More decisively, the ACR is redundant if, across communication-derived task classes, direct mapping into existing workflow standards represents the required action, actor, temporal, conditional, dependency, provenance, and uncertainty semantics without ACR-specific distinctions, and standard extraction metrics discriminate operationally unsafe outputs as well as the proposed executable-correctness measures.

9. Conclusion

Healthcare IT has made the record and then the clinical fact computable; we propose patient-specific clinical intent as the next layer. What characterizes it is not generating clinical text but recovering, from ordinary communication, source-grounded verifiable representations of intended action, formalized as the Actionable Clinical Record upstream of existing workflow infrastructure. We offer the framework, the ACR, and its evaluation framework as reusable constructs to operationalize, evaluate, extend, or falsify.

Data Availability

This article is a conceptual Perspective; no new data were generated or analyzed, and it introduces no new datasets. The illustrative feasibility results summarized in Section 7 are drawn from a companion study;^[52] the controlled benchmark and evaluation described there, including its own data-availability terms, are reported in that study, and no data were newly collected for the present article. The constructs proposed here, the

Actionable Clinical Record schema (Section 5) and the executable-correctness evaluation framework (Section 7), are fully specified in the text and require no separate data deposit. No custom code was produced for this article.

Author Contributions (CRediT)

A. Apartsin: Conceptualization, Methodology, Writing – original draft. Y. Aperstein: Conceptualization, Supervision, Writing – review and editing.

Funding

None declared.

Conflicts of Interest

None declared.

References

1. Ammenwerth E, et al. Twenty-Five Years of Evolution and Hurdles in Electronic Health Records and Interoperability in Medical Research: Comprehensive Review. *J Med Internet Res* 2025;27:e59024. doi:10.2196/59024.
2. Evans RS. Electronic Health Records: Then, Now, and in the Future. *Yearb Med Inform* 2016;25(Suppl 1):S48–S61. doi:10.15265/IYS-2016-soo6.
3. Callen JL, Westbrook JI, Georgiou A, Li J. Failure to follow-up test results for ambulatory patients: a systematic review. *J Gen Intern Med* 2012;27(10):1334–1348. doi:10.1007/s11606-011-1949-5.
4. Patel MP, Schettini P, O'Leary CP, Bosworth HB, Anderson JB, Shah KP. Closing the Referral Loop: an Analysis of Primary Care Referrals to Specialists in a Large Health System. *J Gen Intern Med* 2018;33(5):715–721. doi:10.1007/s11606-018-4392-z.
5. Kwan JL, Singh H. Assigning responsibility to close the loop on radiology test results. *Diagnosis (Berl)* 2017;4(3):173–177. doi:10.1515/dx-2017-0019.
6. Singh H, Meyer AND, Thomas EJ. The frequency of diagnostic errors in outpatient care. *BMJ Qual Saf* 2014;23(9):727–731. doi:10.1136/bmjqs-2013-002627.
7. Slawomirski L, Kelly D, de Bienassis K, Kallas K-A, Klazinga N. The economics of diagnostic safety. *OECD Health Working Papers* 2025. doi:10.1787/fc61057a-en.
8. Wright B, Lennox A, Graber ML, Bragge P. Closing the loop on test results to reduce communication failures: a rapid review of evidence, practice and patient perspectives. *BMC Health Serv Res* 2020;20:897. doi:10.1186/s12913-020-05737-x.

9. Duan-Porter W, Ullman K, Majeski B, Miake-Lye I, Diem S, Wilt TJ. Care Coordination Models and Tools: Systematic Review and Key Informant Interviews. *J Gen Intern Med* 2022;37(6):1367–1379. doi:10.1007/s11606-021-07158-w.
10. Rowley J. The wisdom hierarchy: representations of the DIKW hierarchy. *J Inf Sci* 2007;33(2):163–180. doi:10.1177/0165551506070706.
11. Li R, Niu Y, Scott SR, et al. Using Electronic Medical Record Data for Research in a HIMSS Analytics EMRAM Stage 7 Hospital in Beijing: Cross-sectional Study. *JMIR Med Inform* 2021;9(8):e24405. doi:10.2196/24405.
12. Duncan R, Eden R, Woods L, Wong I, Sullivan C. Synthesizing Dimensions of Digital Maturity in Hospitals: Systematic Review. *J Med Internet Res* 2022;24(3):e32994. doi:10.2196/32994.
13. Blumenthal D, Tavenner M. The “Meaningful Use” Regulation for Electronic Health Records. *N Engl J Med* 2010;363:501–504. doi:10.1056/NEJMp1006114.
14. Adler-Milstein J, et al. Electronic Health Record Adoption In US Hospitals: Progress Continues, But Challenges Persist. *Health Aff* 2015;34(12):2174–2180. doi:10.1377/hlthaff.2015.0992.
15. Weed LL. Medical Records That Guide and Teach. *N Engl J Med* 1968;278:593–600. doi:10.1056/NEJM196803212781204.
16. Bodenreider O, Cornet R, Vreeman DJ. Recent Developments in Clinical Terminologies: SNOMED CT, LOINC, and RxNorm. *Yearb Med Inform* 2018;27(1):129–139. doi:10.1055/s-0038-1667077.
17. Bender D, Sartipi K. HL7 FHIR: An Agile and RESTful Approach to Healthcare Information Exchange. *IEEE CBMS* 2013:326–331. doi:10.1109/CBMS.2013.6627810.
18. Mandel JC, Kreda DA, Mandl KD, Kohane IS, Ramoni RB. SMART on FHIR: a standards-based, interoperable apps platform for electronic health records. *J Am Med Inform Assoc* 2016;23(5):899–908. doi:10.1093/jamia/ocv189.
19. Friedman C, Alderson PO, Austin JHM, Cimino JJ, Johnson SB. A general natural-language text processor for clinical radiology (MedLEE). *J Am Med Inform Assoc* 1994;1(2):161–174. doi:10.1136/jamia.1994.95236146.
20. Savova GK, et al. Mayo clinical Text Analysis and Knowledge Extraction System (cTAKES). *J Am Med Inform Assoc* 2010;17(5):507–513. doi:10.1136/jamia.2009.001560.
21. Safran C, Bloomrosen M, Hammond WE, et al. Toward a National Framework for the Secondary Use of Health Data: An AMIA White Paper. *J Am Med Inform Assoc* 2007;14(1):1–9. doi:10.1197/jamia.M2273.
22. Rosenbloom ST, Denny JC, Xu H, et al. Data from clinical notes: a perspective on the tension between structure and flexible documentation. *J Am Med Inform Assoc* 2011;18(2):181–186. doi:10.1136/jamia.2010.007237.
23. Peleg M. Computer-interpretable clinical guidelines: a methodological review. *J Biomed Inform* 2013;46(4):744–763. doi:10.1016/j.jbi.2013.06.009.
24. Lenz R, Reichert M. IT support for healthcare processes – premises, challenges, perspectives. *Data Knowl Eng* 2007;61(1):39–58. doi:10.1016/j.datak.2006.04.007.
25. Aspland E, Gartner D, Harper P. Clinical pathway modelling: a literature review. *Health Syst* 2021;10(1):1–23. doi:10.1080/20476965.2019.1652547.
26. van der Aalst W. *Process Mining: Data Science in Action*. 2nd ed. Berlin: Springer; 2016. doi:10.1007/978-3-662-49851-4.
27. Rojas E, Munoz-Gama J, Sepúlveda M, Capurro D. Process mining in healthcare: A literature review. *J Biomed Inform* 2016;61:224–236. doi:10.1016/j.jbi.2016.04.007.

28. HL7 FHIR Release 5: Workflow. Health Level Seven International; 2023. <https://hl7.org/fhir/R5/workflow.html> (accessed 9 Aug 2026).
29. Kanaparthi NS, Bakare T, Diao Z, et al. Real-World Evidence Synthesis of Digital Scribes Using Ambient Listening and Generative Artificial Intelligence for Clinician Documentation Workflows: Rapid Review. *JMIR AI* 2025;4:e76743. doi:10.2196/76743.
30. Ben Abacha A, et al. Overview of the MEDIQA-OE 2025 Shared Task on Medical Order Extraction from Doctor-Patient Consultations. *arXiv*; 2025. arXiv:2510.26974.
31. Soysal E, et al. CLAMP: a toolkit for efficiently building customized clinical NLP pipelines. *J Am Med Inform Assoc* 2018;25(3):331–336. doi:10.1093/jamia/ocx132.
32. Yetisgen-Yildiz M, Gunn ML, Xia F, Payne TH. A text processing pipeline to extract recommendations from radiology reports. *J Biomed Inform* 2013;46(2):354–362. doi:10.1016/j.jbi.2012.12.005.
33. Reichenpfader D, Müller H, Denecke K. A scoping review of large language model based approaches for information extraction from radiology reports. *npj Digit Med* 2024;7:222. doi:10.1038/s41746-024-01219-0.
34. Henry S, Buchan K, Filannino M, Stubbs A, Uzuner O. 2018 n2c2 shared task on adverse drug events and medication extraction in electronic health records. *J Am Med Inform Assoc* 2020;27(1):3–12. doi:10.1093/jamia/ocz166.
35. Agrawal M, Heggelmann S, Lang H, Kim Y, Sontag D. Large language models are few-shot clinical information extractors. *Proc. EMNLP* 2022:1998–2022. doi:10.18653/v1/2022.emnlp-main.130.
36. Bethard S, Savova G, Palmer M, Pustejovsky J. SemEval-2017 Task 12: Clinical TempEval. *Proc. 11th Int Workshop Semantic Eval (SemEval-2017)* 2017:565–572. doi:10.18653/v1/S17-2093.
37. Mullenbach J, Pruksachatkun Y, Adler S, et al. CLIP: A Dataset for Extracting Action Items for Physicians from Hospital Discharge Notes. *Proc. 59th Annu Meet Assoc Comput Linguist (ACL-IJCNLP)* 2021:1365–1378. doi:10.18653/v1/2021.acl-long.109.
38. Elgaar M, Cheng J, Vakili N, Amiri H, Celi LA. MedDec: A Dataset for Extracting Medical Decisions from Discharge Summaries. *Findings Assoc Comput Linguist: ACL 2024* 2024:16442–16455. doi:10.18653/v1/2024.findings-acl.975.
39. Nori H, et al. Capabilities of GPT-4 on Medical Challenge Problems. *arXiv* 2023:2303.13375.
40. Ji Z, et al. Survey of Hallucination in Natural Language Generation. *ACM Comput Surv* 2023;55(12):248. doi:10.1145/3571730.
41. Chu Z, Chen J, Chen Q, et al. TimeBench: A Comprehensive Evaluation of Temporal Reasoning Abilities in Large Language Models. *Proc. ACL* 2024. arXiv:2311.17667.
42. Garcez A, Lamb LC. Neurosymbolic AI: the 3rd wave. *Artif Intell Rev* 2023;56:12387–12406. doi:10.1007/s10462-023-10448-w.
43. Prenosil GA, Weitzel TK, Bello SC, et al. Neuro-symbolic AI for auditable cognitive information extraction from medical reports. *Commun Med* 2025;5(1):491. doi:10.1038/s43856-025-01194-x.
44. Uzuner O, South BR, Shen S, DuVall SL. 2010 i2b2/VA challenge on concepts, assertions, and relations in clinical text. *J Am Med Inform Assoc* 2011;18(5):552–556. doi:10.1136/amiajnl-2011-000203.
45. Sun W, Rumshisky A, Uzuner O. Evaluating temporal relations in clinical text: 2012 i2b2 Challenge. *J Am Med Inform Assoc* 2013;20(5):806–813. doi:10.1136/amiajnl-2013-001628.
46. Strötgen J, Gertz M. Multilingual and cross-domain temporal tagging (HeidelTime). *Lang Resour Eval* 2013;47(2):269–298. doi:10.1007/s10579-012-9179-y.
47. Chang AX, Manning CD. SUTime: A library for recognizing and normalizing time expressions. *Proc. LREC* 2012:3735–3740.

48. Knez T, Žitnik S. Multimodal learning for temporal relation extraction in clinical texts. *J Am Med Inform Assoc* 2024;31(6):1380–1387. doi:10.1093/jamia/ocae059.
49. Zolensky A, Jang KJ, Sabin J, et al. Speaker Role Identification in Clinical Conversations. *Pac Symp Biocomput* 2026;31:144–157. doi:10.1101/2025.08.14.25332837.
50. Goddard K, Roudsari A, Wyatt JC. Automation bias: a systematic review of frequency, effect mediators, and mitigators. *J Am Med Inform Assoc* 2012;19(1):121–127. doi:10.1136/amiajnl-2011-000089.
51. Aperstein Y, Gottlib A, Benita G, Apartsin A. Explainable Semantic Text Relations: A Question-Answering Framework for Comparing Document Content. *Information* 2025;16(12):1090. doi:10.3390/info16121090.
52. Laufer M, Aperstein Y, Apartsin A. Reliable Extraction of Clinical Follow-Up Instructions: A Hybrid Neural-Symbolic Pipeline. *arXiv*; 2026. arXiv:2605.26560.
53. Van Tiem J, Cramer E, Iverson C, et al. Listening to the note: clinician perspectives on ambient artificial intelligence scribes in medical documentation. *J Am Med Inform Assoc* 2026;33(2):255–262. doi:10.1093/jamia/ocaf214.
54. Artsi Y, Sorin V, Glicksberg BS, Korfiatis P, Nadkarni GN, Klang E. Large language models in real-world clinical workflows: a systematic review of applications and implementation. *Front Digit Health* 2025;7:1659134. doi:10.3389/fdgth.2025.1659134.